

# hw5q1

April 4, 2025

```
[1]: #sequence length 10

import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
import time
import math
from sklearn.model_selection import train_test_split
import torchinfo
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'Using device: {device}')

# Sample text
```

```

text = '''Next character prediction is a fundamental task in the field of
↳natural language processing (NLP) that involves predicting the next
↳character in a sequence of text based on the characters that precede it.
↳This task is essential for various applications, including text
↳auto-completion, spell checking, and even in the development of
↳sophisticated AI models capable of generating human-like text. At its core,
↳next character prediction relies on statistical models or deep learning
↳algorithms to analyze a given sequence of text and predict which character
↳is most likely to follow. These predictions are based on patterns and
↳relationships learned from large datasets of text during the training phase
↳of the model. One of the most popular approaches to next character
↳prediction involves the use of Recurrent Neural Networks (RNNs), and more
↳specifically, a variant called Long Short-Term Memory (LSTM) networks. RNNs
↳are particularly well-suited for sequential data like text, as they can
↳maintain information in 'memory' about previous characters to inform the
↳prediction of the next character. LSTM networks enhance this capability by
↳being able to remember long-term dependencies, making them even more
↳effective for next character prediction tasks. Training a model for next
↳character prediction involves feeding it large amounts of text data,
↳allowing it to learn the probability of each character's appearance
↳following a sequence of characters. During this training process, the model
↳adjusts its parameters to minimize the difference between its predictions
↳and the actual outcomes, thus improving its predictive accuracy over time.
↳Once trained, the model can be used to predict the next character in a given
↳piece of text by considering the sequence of characters that precede it.
↳This can enhance user experience in text editing software, improve
↳efficiency in coding environments with auto-completion features, and enable
↳more natural interactions with AI-based chatbots and virtual assistants. In
↳summary, next character prediction plays a crucial role in enhancing the
↳capabilities of various NLP applications, making text-based interactions
↳more efficient, accurate, and human-like. Through the use of advanced
↳machine learning models like RNNs and LSTMs, next character prediction
↳continues to evolve, opening new possibilities for the future of text-based
↳technology.'''

```

```

# Preparing the dataset for sequence prediction
max_length = 10 # Maximum length of input sequences
sequences = [text[i:i + max_length] for i in range(len(text) - max_length)]
labels = [text[i + max_length] for i in range(len(text) - max_length)]

# Creating character vocabulary
chars = sorted(list(set(text)))
char_to_ix = {ch: i for i, ch in enumerate(chars)}

# Convert sequences and labels to tensors
X = torch.tensor([[char_to_ix[ch] for ch in seq] for seq in sequences],
↳dtype=torch.long).to(device)

```

```

y = torch.tensor([char_to_ix[label] for label in labels], dtype=torch.long).
    ↪to(device)

# Splitting the dataset into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Define Transformer model
class CharTransformer(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, nhead):
        super(CharTransformer, self).__init__()
        self.embedding = nn.Embedding(input_size, hidden_size)
        encoder_layers = nn.TransformerEncoderLayer(hidden_size, nhead)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers,
    ↪num_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        embedded = self.embedding(x)
        transformer_output = self.transformer_encoder(embedded)
        output = self.fc(transformer_output[:, -1, :]) # Get output of last
    ↪Transformer block
        return output

# Hyperparameters
hidden_size = 128
num_layers = 3
nhead = 2
learning_rate = 0.001
epochs = 50

# Model, loss, and optimizer
model = CharTransformer(len(chars), hidden_size, len(chars), num_layers, nhead).
    ↪to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Display model summary
summary = torchinfo.summary(model, input_data=X_train)
print(summary)

# Training the model
total_start_time = time.time()
for epoch in range(epochs):
    start_time = time.time()
    model.train()
    optimizer.zero_grad()

```

```

output = model(X_train)
loss = criterion(output, y_train)
loss.backward()
optimizer.step()

# Validation
model.eval()
with torch.no_grad():
    val_output = model(X_val)
    val_loss = criterion(val_output, y_val)
    _, predicted = torch.max(val_output, 1)
    val_accuracy = (predicted == y_val).float().mean()

if (epoch+1) % 5 == 0:
    end_time = time.time()
    execution_time = end_time - start_time
    print(f'Epoch {epoch+1}, Loss: {loss.item()}, Validation Loss:
↪{val_loss.item()}, Validation Accuracy: {val_accuracy.item()}, Execution
↪Time: {execution_time} seconds')

total_end_time = time.time()
total_execution_time = total_end_time - total_start_time
print(f'Total Execution Time: {total_execution_time} seconds')

```

Using device: cuda

```

/home/dman/.venv/master/lib/python3.12/site-
packages/torch/nn/modules/transformer.py:385: UserWarning: enable_nested_tensor
is True, but self.use_nested_tensor is False because
encoder_layer.self_attn.batch_first was not True(use batch_first for better
inference performance)
  warnings.warn(

```

```

=====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
CharTransformer                        [1900, 44]                  --
  Embedding: 1-1                      [1900, 10, 128]            5,632
  TransformerEncoder: 1-2             [1900, 10, 128]            --
    ModuleList: 2-1                  --                          --
      TransformerEncoderLayer: 3-1    [1900, 10, 128]            593,024
      TransformerEncoderLayer: 3-2    [1900, 10, 128]            593,024
      TransformerEncoderLayer: 3-3    [1900, 10, 128]            593,024
  Linear: 1-3                        [1900, 44]                 5,676
=====
=====
Total params: 1,790,380

```

```

Trainable params: 1,790,380
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 3.03
=====
=====
Input size (MB): 0.15
Forward/backward pass size (MB): 1129.12
Params size (MB): 6.37
Estimated Total Size (MB): 1135.64
=====
=====
Epoch 5, Loss: 2.8971612453460693, Validation Loss: 2.762009859085083,
Validation Accuracy: 0.21848741173744202, Execution Time: 0.08399343490600586
seconds
Epoch 10, Loss: 2.608732223510742, Validation Loss: 2.5518441200256348,
Validation Accuracy: 0.2920168340206146, Execution Time: 0.13433384895324707
seconds
Epoch 15, Loss: 2.476253032684326, Validation Loss: 2.4380640983581543,
Validation Accuracy: 0.2899159789085388, Execution Time: 0.12561368942260742
seconds
Epoch 20, Loss: 2.403218984603882, Validation Loss: 2.411163330078125,
Validation Accuracy: 0.28151261806488037, Execution Time: 0.1250901222229004
seconds
Epoch 25, Loss: 2.3485896587371826, Validation Loss: 2.371657371520996,
Validation Accuracy: 0.2878151535987854, Execution Time: 0.12447786331176758
seconds
Epoch 30, Loss: 2.3031558990478516, Validation Loss: 2.3365249633789062,
Validation Accuracy: 0.2878151535987854, Execution Time: 0.12671923637390137
seconds
Epoch 35, Loss: 2.279742956161499, Validation Loss: 2.321427583694458,
Validation Accuracy: 0.2899159789085388, Execution Time: 0.12458181381225586
seconds
Epoch 40, Loss: 2.2605137825012207, Validation Loss: 2.316436529159546,
Validation Accuracy: 0.2857142984867096, Execution Time: 0.12621021270751953
seconds
Epoch 45, Loss: 2.2435638904571533, Validation Loss: 2.311800956726074,
Validation Accuracy: 0.2878151535987854, Execution Time: 0.12342119216918945
seconds
Epoch 50, Loss: 2.237259864807129, Validation Loss: 2.3116962909698486,
Validation Accuracy: 0.2878151535987854, Execution Time: 0.12564873695373535
seconds
Total Execution Time: 6.421080589294434 seconds

```

```

[2]: #sequence length 20
      # Sample text

```

```

text = '''Next character prediction is a fundamental task in the field of
↳natural language processing (NLP) that involves predicting the next
↳character in a sequence of text based on the characters that precede it.
↳This task is essential for various applications, including text
↳auto-completion, spell checking, and even in the development of
↳sophisticated AI models capable of generating human-like text. At its core,
↳next character prediction relies on statistical models or deep learning
↳algorithms to analyze a given sequence of text and predict which character
↳is most likely to follow. These predictions are based on patterns and
↳relationships learned from large datasets of text during the training phase
↳of the model. One of the most popular approaches to next character
↳prediction involves the use of Recurrent Neural Networks (RNNs), and more
↳specifically, a variant called Long Short-Term Memory (LSTM) networks. RNNs
↳are particularly well-suited for sequential data like text, as they can
↳maintain information in 'memory' about previous characters to inform the
↳prediction of the next character. LSTM networks enhance this capability by
↳being able to remember long-term dependencies, making them even more
↳effective for next character prediction tasks. Training a model for next
↳character prediction involves feeding it large amounts of text data,
↳allowing it to learn the probability of each character's appearance
↳following a sequence of characters. During this training process, the model
↳adjusts its parameters to minimize the difference between its predictions
↳and the actual outcomes, thus improving its predictive accuracy over time.
↳Once trained, the model can be used to predict the next character in a given
↳piece of text by considering the sequence of characters that precede it.
↳This can enhance user experience in text editing software, improve
↳efficiency in coding environments with auto-completion features, and enable
↳more natural interactions with AI-based chatbots and virtual assistants. In
↳summary, next character prediction plays a crucial role in enhancing the
↳capabilities of various NLP applications, making text-based interactions
↳more efficient, accurate, and human-like. Through the use of advanced
↳machine learning models like RNNs and LSTMs, next character prediction
↳continues to evolve, opening new possibilities for the future of text-based
↳technology.'''

```

```

# Preparing the dataset for sequence prediction
max_length = 20 # Maximum length of input sequences
sequences = [text[i:i + max_length] for i in range(len(text) - max_length)]
labels = [text[i + max_length] for i in range(len(text) - max_length)]

# Creating character vocabulary
chars = sorted(list(set(text)))
char_to_ix = {ch: i for i, ch in enumerate(chars)}

# Convert sequences and labels to tensors
X = torch.tensor([[char_to_ix[ch] for ch in seq] for seq in sequences],
↳dtype=torch.long).to(device)

```

```

y = torch.tensor([char_to_ix[label] for label in labels], dtype=torch.long).
    ↪to(device)

# Splitting the dataset into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Define Transformer model
class CharTransformer(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, nhead):
        super(CharTransformer, self).__init__()
        self.embedding = nn.Embedding(input_size, hidden_size)
        encoder_layers = nn.TransformerEncoderLayer(hidden_size, nhead)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers,
    ↪num_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        embedded = self.embedding(x)
        transformer_output = self.transformer_encoder(embedded)
        output = self.fc(transformer_output[:, -1, :]) # Get output of last
    ↪Transformer block
        return output

# Hyperparameters
hidden_size = 128
num_layers = 3
nhead = 2
learning_rate = 0.001
epochs = 50

# Model, loss, and optimizer
model = CharTransformer(len(chars), hidden_size, len(chars), num_layers, nhead).
    ↪to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Display model summary
summary = torchinfo.summary(model, input_data=X_train)
print(summary)

# Training the model
total_start_time = time.time()
for epoch in range(epochs):
    start_time = time.time()
    model.train()
    optimizer.zero_grad()

```

```

output = model(X_train)
loss = criterion(output, y_train)
loss.backward()
optimizer.step()

# Validation
model.eval()
with torch.no_grad():
    val_output = model(X_val)
    val_loss = criterion(val_output, y_val)
    _, predicted = torch.max(val_output, 1)
    val_accuracy = (predicted == y_val).float().mean()

if (epoch+1) % 5 == 0:
    end_time = time.time()
    execution_time = end_time - start_time
    print(f'Epoch {epoch+1}, Loss: {loss.item()}, Validation Loss:
↪{val_loss.item()}, Validation Accuracy: {val_accuracy.item()}, Execution
↪Time: {execution_time} seconds')

total_end_time = time.time()
total_execution_time = total_end_time - total_start_time
print(f'Total Execution Time: {total_execution_time} seconds')

```

```

=====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
CharTransformer                       [1892, 44]                  --
  Embedding: 1-1                      [1892, 20, 128]            5,632
  TransformerEncoder: 1-2             [1892, 20, 128]            --
    ModuleList: 2-1                  --                          --
      TransformerEncoderLayer: 3-1    [1892, 20, 128]            593,024
      TransformerEncoderLayer: 3-2    [1892, 20, 128]            593,024
      TransformerEncoderLayer: 3-3    [1892, 20, 128]            593,024
  Linear: 1-3                         [1892, 44]                 5,676
=====
=====
Total params: 1,790,380
Trainable params: 1,790,380
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 3.01
=====
=====
Input size (MB): 0.30
Forward/backward pass size (MB): 2248.06
Params size (MB): 6.37

```



Estimated Total Size (MB): 2254.73

=====

```
/home/dman/.venv/master/lib/python3.12/site-  
packages/torch/nn/modules/transformer.py:385: UserWarning: enable_nested_tensor  
is True, but self.use_nested_tensor is False because  
encoder_layer.self_attn.batch_first was not True(use batch_first for better  
inference performance)  
  warnings.warn(
```

Epoch 5, Loss: 2.8855316638946533, Validation Loss: 2.79542875289917, Validation  
Accuracy: 0.2109704613685608, Execution Time: 0.21418094635009766 seconds

Epoch 10, Loss: 2.5777204036712646, Validation Loss: 2.5451254844665527,  
Validation Accuracy: 0.23839661478996277, Execution Time: 0.24049806594848633  
seconds

Epoch 15, Loss: 2.4539241790771484, Validation Loss: 2.4622461795806885,  
Validation Accuracy: 0.2805907130241394, Execution Time: 0.24182915687561035  
seconds

Epoch 20, Loss: 2.3868677616119385, Validation Loss: 2.4474551677703857,  
Validation Accuracy: 0.2763713002204895, Execution Time: 0.24112629890441895  
seconds

Epoch 25, Loss: 2.3527097702026367, Validation Loss: 2.416934013366699,  
Validation Accuracy: 0.24894513189792633, Execution Time: 0.24055862426757812  
seconds

Epoch 30, Loss: 2.310042142868042, Validation Loss: 2.386185646057129,  
Validation Accuracy: 0.25738394260406494, Execution Time: 0.2419743537902832  
seconds

Epoch 35, Loss: 2.2733840942382812, Validation Loss: 2.382976531982422,  
Validation Accuracy: 0.25316452980041504, Execution Time: 0.2410719394683838  
seconds

Epoch 40, Loss: 2.264220714569092, Validation Loss: 2.361839532852173,  
Validation Accuracy: 0.25316452980041504, Execution Time: 0.24018406867980957  
seconds

Epoch 45, Loss: 2.2494311332702637, Validation Loss: 2.3602750301361084,  
Validation Accuracy: 0.2679324746131897, Execution Time: 0.23980093002319336  
seconds

Epoch 50, Loss: 2.2300937175750732, Validation Loss: 2.364194869995117,  
Validation Accuracy: 0.27426159381866455, Execution Time: 0.24025893211364746  
seconds

Total Execution Time: 11.98580002784729 seconds

[3]: *#sequence length 30*

*# Sample text*

```

text = '''Next character prediction is a fundamental task in the field of
↳natural language processing (NLP) that involves predicting the next
↳character in a sequence of text based on the characters that precede it.
↳This task is essential for various applications, including text
↳auto-completion, spell checking, and even in the development of
↳sophisticated AI models capable of generating human-like text. At its core,
↳next character prediction relies on statistical models or deep learning
↳algorithms to analyze a given sequence of text and predict which character
↳is most likely to follow. These predictions are based on patterns and
↳relationships learned from large datasets of text during the training phase
↳of the model. One of the most popular approaches to next character
↳prediction involves the use of Recurrent Neural Networks (RNNs), and more
↳specifically, a variant called Long Short-Term Memory (LSTM) networks. RNNs
↳are particularly well-suited for sequential data like text, as they can
↳maintain information in 'memory' about previous characters to inform the
↳prediction of the next character. LSTM networks enhance this capability by
↳being able to remember long-term dependencies, making them even more
↳effective for next character prediction tasks. Training a model for next
↳character prediction involves feeding it large amounts of text data,
↳allowing it to learn the probability of each character's appearance
↳following a sequence of characters. During this training process, the model
↳adjusts its parameters to minimize the difference between its predictions
↳and the actual outcomes, thus improving its predictive accuracy over time.
↳Once trained, the model can be used to predict the next character in a given
↳piece of text by considering the sequence of characters that precede it.
↳This can enhance user experience in text editing software, improve
↳efficiency in coding environments with auto-completion features, and enable
↳more natural interactions with AI-based chatbots and virtual assistants. In
↳summary, next character prediction plays a crucial role in enhancing the
↳capabilities of various NLP applications, making text-based interactions
↳more efficient, accurate, and human-like. Through the use of advanced
↳machine learning models like RNNs and LSTMs, next character prediction
↳continues to evolve, opening new possibilities for the future of text-based
↳technology.'''

```

```

# Preparing the dataset for sequence prediction
max_length = 30 # Maximum length of input sequences
sequences = [text[i:i + max_length] for i in range(len(text) - max_length)]
labels = [text[i + max_length] for i in range(len(text) - max_length)]

# Creating character vocabulary
chars = sorted(list(set(text)))
char_to_ix = {ch: i for i, ch in enumerate(chars)}

# Convert sequences and labels to tensors
X = torch.tensor([[char_to_ix[ch] for ch in seq] for seq in sequences],
↳dtype=torch.long).to(device)

```

```

y = torch.tensor([char_to_ix[label] for label in labels], dtype=torch.long).
    ↪to(device)

# Splitting the dataset into training and validation sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Define Transformer model
class CharTransformer(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, nhead):
        super(CharTransformer, self).__init__()
        self.embedding = nn.Embedding(input_size, hidden_size)
        encoder_layers = nn.TransformerEncoderLayer(hidden_size, nhead)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers,
    ↪num_layers)
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        embedded = self.embedding(x)
        transformer_output = self.transformer_encoder(embedded)
        output = self.fc(transformer_output[:, -1, :]) # Get output of last
    ↪Transformer block
        return output

# Hyperparameters
hidden_size = 128
num_layers = 3
nhead = 2
learning_rate = 0.001
epochs = 50

# Model, loss, and optimizer
model = CharTransformer(len(chars), hidden_size, len(chars), num_layers, nhead).
    ↪to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Display model summary
summary = torchinfo.summary(model, input_data=X_train)
print(summary)

# Training the model
total_start_time = time.time()
for epoch in range(epochs):
    start_time = time.time()
    model.train()
    optimizer.zero_grad()

```

```

output = model(X_train)
loss = criterion(output, y_train)
loss.backward()
optimizer.step()

# Validation
model.eval()
with torch.no_grad():
    val_output = model(X_val)
    val_loss = criterion(val_output, y_val)
    _, predicted = torch.max(val_output, 1)
    val_accuracy = (predicted == y_val).float().mean()

if (epoch+1) % 5 == 0:
    end_time = time.time()
    execution_time = end_time - start_time
    print(f'Epoch {epoch+1}, Loss: {loss.item()}, Validation Loss:
↪{val_loss.item()}, Validation Accuracy: {val_accuracy.item()}, Execution
↪Time: {execution_time} seconds')

total_end_time = time.time()
total_execution_time = total_end_time - total_start_time
print(f'Total Execution Time: {total_execution_time} seconds')

```

```

=====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
CharTransformer                        [1884, 44]                  --
  Embedding: 1-1                       [1884, 30, 128]             5,632
  TransformerEncoder: 1-2               [1884, 30, 128]             --
    ModuleList: 2-1
      TransformerEncoderLayer: 3-1       [1884, 30, 128]             593,024
      TransformerEncoderLayer: 3-2       [1884, 30, 128]             593,024
      TransformerEncoderLayer: 3-3       [1884, 30, 128]             593,024
  Linear: 1-3                           [1884, 44]                  5,676
=====
=====
Total params: 1,790,380
Trainable params: 1,790,380
Non-trainable params: 0
Total mult-adds (Units.GIGABYTES): 3.00
=====
=====
Input size (MB): 0.45
Forward/backward pass size (MB): 3357.50
Params size (MB): 6.37

```

Estimated Total Size (MB): 3364.32

=====

Epoch 5, Loss: 2.9286999702453613, Validation Loss: 2.8957574367523193,  
Validation Accuracy: 0.17796610295772552, Execution Time: 0.3450007438659668  
seconds

Epoch 10, Loss: 2.6100516319274902, Validation Loss: 2.6570611000061035,  
Validation Accuracy: 0.24152542650699615, Execution Time: 0.3161318302154541  
seconds

Epoch 15, Loss: 2.4706196784973145, Validation Loss: 2.5679800510406494,  
Validation Accuracy: 0.2330508530139923, Execution Time: 0.2823753356933594  
seconds

Epoch 20, Loss: 2.394347906112671, Validation Loss: 2.5069236755371094,  
Validation Accuracy: 0.24576270580291748, Execution Time: 0.28311777114868164  
seconds

Epoch 25, Loss: 2.3387489318847656, Validation Loss: 2.484586477279663,  
Validation Accuracy: 0.24788135290145874, Execution Time: 0.2780444622039795  
seconds

Epoch 30, Loss: 2.2942371368408203, Validation Loss: 2.457167625427246,  
Validation Accuracy: 0.24576270580291748, Execution Time: 0.28285980224609375  
seconds

Epoch 35, Loss: 2.26870059967041, Validation Loss: 2.449058771133423, Validation  
Accuracy: 0.25, Execution Time: 0.28238797187805176 seconds

Epoch 40, Loss: 2.2444820404052734, Validation Loss: 2.4435513019561768,  
Validation Accuracy: 0.24788135290145874, Execution Time: 0.31443238258361816  
seconds

Epoch 45, Loss: 2.232463836669922, Validation Loss: 2.4325082302093506,  
Validation Accuracy: 0.24788135290145874, Execution Time: 0.2831428050994873  
seconds

Epoch 50, Loss: 2.2242465019226074, Validation Loss: 2.435275077819824,  
Validation Accuracy: 0.24788135290145874, Execution Time: 0.27874326705932617  
seconds

Total Execution Time: 14.616542339324951 seconds