

# rnn-base

April 4, 2025

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import numpy as np
from torch.utils.data import Dataset, DataLoader
import math
import time
import matplotlib.pyplot as plt

# Use GPU if available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

#Sample text
```

```

text = '''Next character prediction is a fundamental task in the field of
↳natural language processing (NLP) that involves predicting the next
↳character in a sequence of text based on the characters that precede it.
↳This task is essential for various applications, including text
↳auto-completion, spell checking, and even in the development of
↳sophisticated AI models capable of generating human-like text. At its core,
↳next character prediction relies on statistical models or deep learning
↳algorithms to analyze a given sequence of text and predict which character
↳is most likely to follow. These predictions are based on patterns and
↳relationships learned from large datasets of text during the training phase
↳of the model. One of the most popular approaches to next character
↳prediction involves the use of Recurrent Neural Networks (RNNs), and more
↳specifically, a variant called Long Short-Term Memory (LSTM) networks. RNNs
↳are particularly well-suited for sequential data like text, as they can
↳maintain information in 'memory' about previous characters to inform the
↳prediction of the next character. LSTM networks enhance this capability by
↳being able to remember long-term dependencies, making them even more
↳effective for next character prediction tasks. Training a model for next
↳character prediction involves feeding it large amounts of text data,
↳allowing it to learn the probability of each character's appearance
↳following a sequence of characters. During this training process, the model
↳adjusts its parameters to minimize the difference between its predictions
↳and the actual outcomes, thus improving its predictive accuracy over time.
↳Once trained, the model can be used to predict the next character in a given
↳piece of text by considering the sequence of characters that precede it.
↳This can enhance user experience in text editing software, improve
↳efficiency in coding environments with auto-completion features, and enable
↳more natural interactions with AI-based chatbots and virtual assistants. In
↳summary, next character prediction plays a crucial role in enhancing the
↳capabilities of various NLP applications, making text-based interactions
↳more efficient, accurate, and human-like. Through the use of advanced
↳machine learning models like RNNs and LSTMs, next character prediction
↳continues to evolve, opening new possibilities for the future of text-based
↳technology.'''

```

```

# Create character-level vocabulary
chars = sorted(list(set(text)))
vocab_size = len(chars)
print("Unique characters:", vocab_size)
char2idx = {ch: i for i, ch in enumerate(chars)}
idx2char = {i: ch for i, ch in enumerate(chars)}

# Dataset that produces input sequences and target character (the next
↳character)
class CharDataset(Dataset):
    def __init__(self, text, seq_length=20):

```

```

        self.text = text
        self.seq_length = seq_length
        self.data = [char2idx[ch] for ch in text]

    def __len__(self):
        return len(self.data) - self.seq_length

    def __getitem__(self, idx):
        # Input sequence of length seq_length and target: the next character
        x = torch.tensor(self.data[idx:idx+self.seq_length], dtype=torch.long)
        y = torch.tensor(self.data[idx+self.seq_length], dtype=torch.long)
        return x, y

# Choose sequence length: try 10, 20, or 30 (here, we use 20)
seq_length = 20
dataset = CharDataset(text, seq_length=seq_length)
train_size = int(0.9 * len(dataset))
val_size = len(dataset) - train_size
train_dataset, val_dataset = torch.utils.data.random_split(dataset,
    ↪ [train_size, val_size])
batch_size = 64
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size)

# Model 1: Plain RNN-based model (using LSTM)
class RNNModel(nn.Module):
    def __init__(self, vocab_size, embed_dim, hidden_dim, num_layers=1):
        super(RNNModel, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim)
        self.lstm = nn.LSTM(embed_dim, hidden_dim, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_dim, vocab_size)

    def forward(self, x):
        # x: (batch, seq_length)
        embedded = self.embedding(x)                # (batch, seq_length, ↪
    ↪ embed_dim)
        output, (hn, cn) = self.lstm(embedded)      # output: (batch, ↪
    ↪ seq_length, hidden_dim)
        # Use the final hidden state for prediction
        out = self.fc(hn[-1])                      # (batch, vocab_size)
        return out

# Model 2: RNN with cross attention.
# After processing the sequence with LSTM, we attend over all time steps using ↪
    ↪ the final hidden state as query.
class RNNWithAttention(nn.Module):
    def __init__(self, vocab_size, embed_dim, hidden_dim, num_layers=1):

```

```

    super(RNNWithAttention, self).__init__()
    self.embedding = nn.Embedding(vocab_size, embed_dim)
    self.lstm = nn.LSTM(embed_dim, hidden_dim, num_layers, batch_first=True)
    # Attention parameters
    self.attn_linear = nn.Linear(hidden_dim, hidden_dim)
    self.fc = nn.Linear(hidden_dim * 2, vocab_size)

    def forward(self, x):
        # x: (batch, seq_length)
        embedded = self.embedding(x)          # (batch, seq_length, ↵
        ↵embed_dim)
        outputs, (hn, cn) = self.lstm(embedded) # outputs: (batch, ↵
        ↵seq_length, hidden_dim)
        # Use the final hidden state as the query for attention
        query = hn[-1].unsqueeze(1)          # (batch, 1, hidden_dim)
        # Project outputs to get keys
        keys = self.attn_linear(outputs)      # (batch, seq_length, ↵
        ↵hidden_dim)
        # Compute dot-product attention scores (scaled)
        scores = torch.bmm(query, keys.transpose(1, 2)) / math.sqrt(keys.
        ↵size(-1)) # (batch, 1, seq_length)
        attn_weights = torch.softmax(scores, dim=-1) # (batch, 1, seq_length)
        # Compute context vector as weighted sum of outputs
        context = torch.bmm(attn_weights, outputs) # (batch, 1, hidden_dim)
        context = context.squeeze(1)          # (batch, hidden_dim)
        query = query.squeeze(1)             # (batch, hidden_dim)
        combined = torch.cat((query, context), dim=1) # (batch, 2*hidden_dim)
        out = self.fc(combined)              # (batch, vocab_size)
        return out

# Hyperparameters
embed_dim = 128
hidden_dim = 256
num_layers = 1
learning_rate = 0.003
num_epochs = 20

# Instantiate models
model_plain = RNNModel(vocab_size, embed_dim, hidden_dim, num_layers).to(device)
model_attn = RNNWithAttention(vocab_size, embed_dim, hidden_dim, num_layers).
    ↵to(device)

criterion = nn.CrossEntropyLoss()
optimizer_plain = optim.Adam(model_plain.parameters(), lr=learning_rate)
optimizer_attn = optim.Adam(model_attn.parameters(), lr=learning_rate)

# Functions to train and evaluate a model

```

```

def train_model(model, optimizer, loader):
    model.train()
    total_loss = 0
    for x, y in loader:
        x, y = x.to(device), y.to(device)
        optimizer.zero_grad()
        logits = model(x)          # logits shape: (batch, vocab_size)
        loss = criterion(logits, y)
        loss.backward()
        optimizer.step()
        total_loss += loss.item() * x.size(0)
    return total_loss / len(loader.dataset)

def evaluate_model(model, loader):
    model.eval()
    total_loss = 0
    correct = 0
    with torch.no_grad():
        for x, y in loader:
            x, y = x.to(device), y.to(device)
            logits = model(x)
            loss = criterion(logits, y)
            total_loss += loss.item() * x.size(0)
            preds = logits.argmax(dim=1)
            correct += (preds == y).sum().item()
    avg_loss = total_loss / len(loader.dataset)
    accuracy = correct / len(loader.dataset)
    return avg_loss, accuracy

# Containers for metrics
plain_train_losses = []
plain_val_losses = []
plain_val_accs = []

attn_train_losses = []
attn_val_losses = []
attn_val_accs = []

# Training loop for the plain RNN model
print("Training Plain RNN Model (without attention)")
start_time = time.time()
for epoch in range(num_epochs):
    train_loss = train_model(model_plain, optimizer_plain, train_loader)
    val_loss, val_acc = evaluate_model(model_plain, val_loader)
    plain_train_losses.append(train_loss)
    plain_val_losses.append(val_loss)
    plain_val_accs.append(val_acc)

```

```

        print(f"Plain RNN Epoch {epoch+1}/{num_epochs}: Train Loss: {train_loss:.4f}, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")
    plain_time = time.time() - start_time

    # Training loop for the RNN with attention model
    print("\nTraining RNN Model with Cross Attention")
    start_time = time.time()
    for epoch in range(num_epochs):
        train_loss = train_model(model_attn, optimizer_attn, train_loader)
        val_loss, val_acc = evaluate_model(model_attn, val_loader)
        attn_train_losses.append(train_loss)
        attn_val_losses.append(val_loss)
        attn_val_accs.append(val_acc)
        print(f"Attention RNN Epoch {epoch+1}/{num_epochs}: Train Loss: {train_loss:.4f}, Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.4f}")
    attn_time = time.time() - start_time

    print("\nTiming Summary:")
    print(f"Plain RNN training time: {plain_time:.2f} seconds")
    print(f"RNN with Attention training time: {attn_time:.2f} seconds")

    # Visualization of training loss and validation accuracy
    epochs = range(1, num_epochs+1)
    plt.figure(figsize=(12, 5))

    # Plot training loss
    plt.subplot(1, 2, 1)
    plt.plot(epochs, plain_train_losses, label='Plain RNN Train Loss')
    plt.plot(epochs, attn_train_losses, label='Attention RNN Train Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Training Loss')
    plt.title('Training Loss Comparison')
    plt.legend()
    plt.grid(True)

    # Plot validation accuracy
    plt.subplot(1, 2, 2)
    plt.plot(epochs, plain_val_accs, label='Plain RNN Val Accuracy')
    plt.plot(epochs, attn_val_accs, label='Attention RNN Val Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Validation Accuracy')
    plt.title('Validation Accuracy Comparison')
    plt.legend()
    plt.grid(True)

    plt.tight_layout()
    plt.show()

```

Using device: cuda

Unique characters: 44

Training Plain RNN Model (without attention)

Plain RNN Epoch 1/20: Train Loss: 2.8618, Val Loss: 2.2864, Val Acc: 0.3502  
Plain RNN Epoch 2/20: Train Loss: 2.2006, Val Loss: 1.9352, Val Acc: 0.4852  
Plain RNN Epoch 3/20: Train Loss: 1.8559, Val Loss: 1.7766, Val Acc: 0.5232  
Plain RNN Epoch 4/20: Train Loss: 1.5634, Val Loss: 1.7666, Val Acc: 0.5274  
Plain RNN Epoch 5/20: Train Loss: 1.3367, Val Loss: 1.7145, Val Acc: 0.5359  
Plain RNN Epoch 6/20: Train Loss: 1.1408, Val Loss: 1.6906, Val Acc: 0.5401  
Plain RNN Epoch 7/20: Train Loss: 0.9630, Val Loss: 1.7211, Val Acc: 0.5527  
Plain RNN Epoch 8/20: Train Loss: 0.7948, Val Loss: 1.6874, Val Acc: 0.5949  
Plain RNN Epoch 9/20: Train Loss: 0.6615, Val Loss: 1.7467, Val Acc: 0.5485  
Plain RNN Epoch 10/20: Train Loss: 0.5541, Val Loss: 1.8182, Val Acc: 0.5696  
Plain RNN Epoch 11/20: Train Loss: 0.4583, Val Loss: 1.8762, Val Acc: 0.5781  
Plain RNN Epoch 12/20: Train Loss: 0.3645, Val Loss: 1.9161, Val Acc: 0.5612  
Plain RNN Epoch 13/20: Train Loss: 0.3026, Val Loss: 1.9701, Val Acc: 0.5738  
Plain RNN Epoch 14/20: Train Loss: 0.2499, Val Loss: 2.0183, Val Acc: 0.5654  
Plain RNN Epoch 15/20: Train Loss: 0.2161, Val Loss: 2.1004, Val Acc: 0.5443  
Plain RNN Epoch 16/20: Train Loss: 0.1808, Val Loss: 2.1231, Val Acc: 0.5654  
Plain RNN Epoch 17/20: Train Loss: 0.1460, Val Loss: 2.1459, Val Acc: 0.5738  
Plain RNN Epoch 18/20: Train Loss: 0.1221, Val Loss: 2.2190, Val Acc: 0.5696  
Plain RNN Epoch 19/20: Train Loss: 0.0950, Val Loss: 2.2381, Val Acc: 0.5570  
Plain RNN Epoch 20/20: Train Loss: 0.0914, Val Loss: 2.2758, Val Acc: 0.5570

Training RNN Model with Cross Attention

Attention RNN Epoch 1/20: Train Loss: 2.8564, Val Loss: 2.2420, Val Acc: 0.3586  
Attention RNN Epoch 2/20: Train Loss: 2.1630, Val Loss: 1.9462, Val Acc: 0.4473  
Attention RNN Epoch 3/20: Train Loss: 1.7594, Val Loss: 1.8078, Val Acc: 0.5443  
Attention RNN Epoch 4/20: Train Loss: 1.4078, Val Loss: 1.8091, Val Acc: 0.5401  
Attention RNN Epoch 5/20: Train Loss: 1.1126, Val Loss: 1.7581, Val Acc: 0.5823  
Attention RNN Epoch 6/20: Train Loss: 0.8244, Val Loss: 1.7882, Val Acc: 0.5570  
Attention RNN Epoch 7/20: Train Loss: 0.6112, Val Loss: 1.8913, Val Acc: 0.5527  
Attention RNN Epoch 8/20: Train Loss: 0.4597, Val Loss: 2.0242, Val Acc: 0.5443  
Attention RNN Epoch 9/20: Train Loss: 0.3382, Val Loss: 2.0884, Val Acc: 0.5527  
Attention RNN Epoch 10/20: Train Loss: 0.2238, Val Loss: 2.1856, Val Acc: 0.5612  
Attention RNN Epoch 11/20: Train Loss: 0.1424, Val Loss: 2.2690, Val Acc: 0.5485  
Attention RNN Epoch 12/20: Train Loss: 0.0938, Val Loss: 2.3573, Val Acc: 0.5527  
Attention RNN Epoch 13/20: Train Loss: 0.0656, Val Loss: 2.4867, Val Acc: 0.5443  
Attention RNN Epoch 14/20: Train Loss: 0.0481, Val Loss: 2.5340, Val Acc: 0.5485  
Attention RNN Epoch 15/20: Train Loss: 0.0422, Val Loss: 2.5659, Val Acc: 0.5527  
Attention RNN Epoch 16/20: Train Loss: 0.0381, Val Loss: 2.5819, Val Acc: 0.5527  
Attention RNN Epoch 17/20: Train Loss: 0.0391, Val Loss: 2.6496, Val Acc: 0.5485  
Attention RNN Epoch 18/20: Train Loss: 0.0287, Val Loss: 2.6576, Val Acc: 0.5612  
Attention RNN Epoch 19/20: Train Loss: 0.0272, Val Loss: 2.6998, Val Acc: 0.5485  
Attention RNN Epoch 20/20: Train Loss: 0.0265, Val Loss: 2.7093, Val Acc: 0.5527

Timing Summary:

Plain RNN training time: 1.52 seconds

RNN with Attention training time: 1.60 seconds

