

# homework5

November 15, 2024

[3]: *#ORIGINAL CODE FROM LECTURE NOTES*

```
import torch
import torch.optim as optim

# Training data
t_c = [0.5, 14.0, 15.0, 28.0, 11.0, 8.0, 3.0, -4.0, 6.0, 13.0, 21.0]
t_u = [35.7, 55.9, 58.2, 81.9, 56.3, 48.9, 33.9, 21.8, 48.4, 60.4, 68.4]
t_c = torch.tensor(t_c)
t_u = torch.tensor(t_u)

# Normalize the input
t_un = 0.1 * t_u

def model(t_u, w, b):
    return w * t_u + b

def loss_fn(t_p, t_c):
    squared_diffs = (t_p - t_c)**2
    return squared_diffs.mean()

params = torch.tensor([1.0, 0.0], requires_grad=True)
optimizer = optim.Adam([params], lr=1e-2)

n_epochs = 5000
for epoch in range(n_epochs):
    t_p = model(t_un, *params)
    loss = loss_fn(t_p, t_c)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    if epoch % 500 == 0:
        print(f'Epoch {epoch}, Loss {loss.item():.4f}')

print(f'\nFinal Parameters: w={params[0].item():.4f}, b={params[1].item():.4f}')
```

Epoch 0, Loss 80.3643  
Epoch 500, Loss 24.9258  
Epoch 1000, Loss 15.7372  
Epoch 1500, Loss 9.4454  
Epoch 2000, Loss 5.7623  
Epoch 2500, Loss 3.9305  
Epoch 3000, Loss 3.1960  
Epoch 3500, Loss 2.9770  
Epoch 4000, Loss 2.9332  
Epoch 4500, Loss 2.9280

Final Parameters: w=5.3660, b=-17.2952

```
[7]: ##Problem 1
      #1a and 1b

      import torch
      import torch.optim as optim
      import matplotlib.pyplot as plt

      # Training data
      t_c = [0.5, 14.0, 15.0, 28.0, 11.0, 8.0, 3.0, -4.0, 6.0, 13.0, 21.0]
      t_u = [35.7, 55.9, 58.2, 81.9, 56.3, 48.9, 33.9, 21.8, 48.4, 60.4, 68.4]
      t_c = torch.tensor(t_c, dtype=torch.float)
      t_u = torch.tensor(t_u, dtype=torch.float)

      # Normalize input
      t_un = (t_u - torch.mean(t_u)) / torch.std(t_u)

      def model(t_u, w2, w1, b):
          return w2 * t_u**2 + w1 * t_u + b

      def loss_fn(t_p, t_c):
          squared_diffs = (t_p - t_c)**2
          return squared_diffs.mean()

      learning_rates = [0.1, 0.01, 0.001, 0.0001]
      final_losses = []
      all_params = []

      for lr in learning_rates:
          params = torch.tensor([0.1, 0.1, 0.0], requires_grad=True)
          optimizer = optim.Adam([params], lr=lr)

          print(f"\nTraining with learning rate: {lr}")
          for epoch in range(5000):
              t_p = model(t_un, *params)
```

```

        loss = loss_fn(t_p, t_c)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if epoch % 500 == 0:
        print(f'Epoch {epoch}, Loss {loss.item():.4f}')

    final_losses.append(loss.item())
    all_params.append(params.detach().clone())
    print(f'Final parameters: w2={params[0].item():.4f}, w1={params[1].item():.4f}, b={params[2].item():.4f}')

# Find best model
best_lr_index = final_losses.index(min(final_losses))
best_params = all_params[best_lr_index]

# Plot results
plt.figure(figsize=(10, 6))
t_u_range = torch.linspace(min(t_un), max(t_un), 100)
predictions = model(t_u_range, *best_params)

plt.scatter(t_un.numpy(), t_c.numpy(), label='Data')
plt.plot(t_u_range.numpy(), predictions.detach().numpy(), 'r-',
        label='Nonlinear Model')
plt.xlabel('Normalized Input Temperature')
plt.ylabel('Output Temperature')
plt.legend()
plt.title('Nonlinear Temperature Prediction Model')
plt.grid(True)
plt.show()

print(f"\nBest learning rate: {learning_rates[best_lr_index]}")
print(f"Best final loss: {min(final_losses):.4f}")

```

```

Training with learning rate: 0.1
Epoch 0, Loss 183.7949
Epoch 500, Loss 2.0907
Epoch 1000, Loss 2.0907
Epoch 1500, Loss 2.0907
Epoch 2000, Loss 2.0907
Epoch 2500, Loss 2.0907
Epoch 3000, Loss 2.0907
Epoch 3500, Loss 2.0907
Epoch 4000, Loss 2.0907

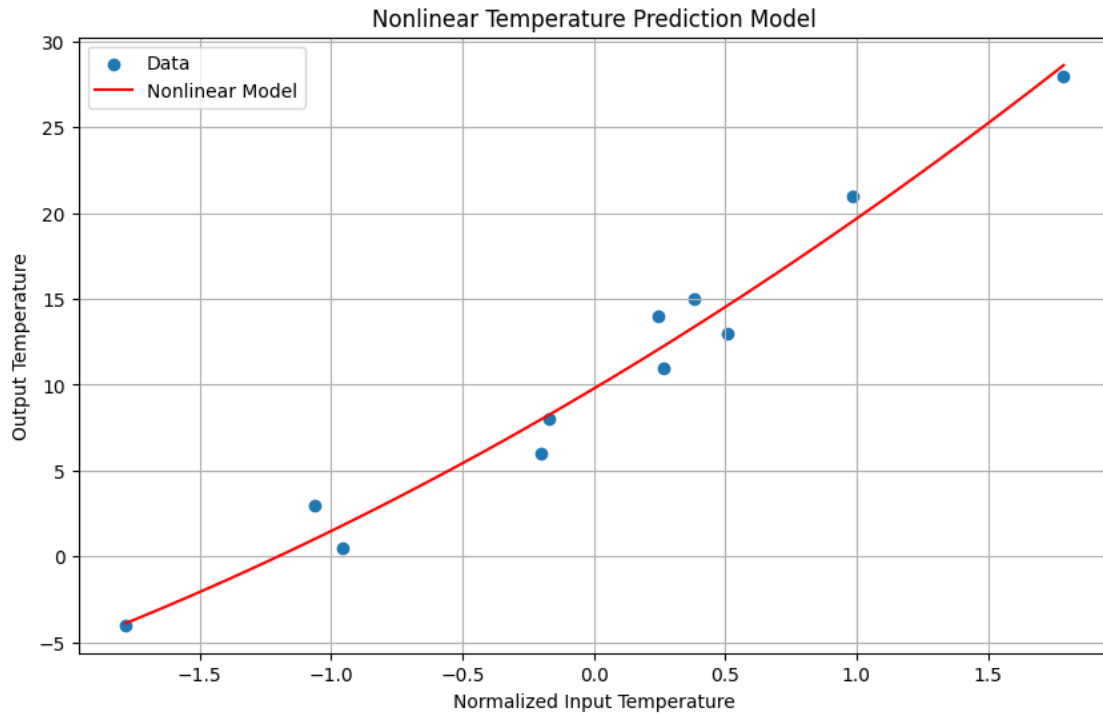
```

Epoch 4500, Loss 2.0907  
Final parameters: w2=0.8019, w1=9.1033, b=9.7710

Training with learning rate: 0.01  
Epoch 0, Loss 183.7949  
Epoch 500, Loss 44.1960  
Epoch 1000, Loss 12.4033  
Epoch 1500, Loss 3.9486  
Epoch 2000, Loss 2.3267  
Epoch 2500, Loss 2.1083  
Epoch 3000, Loss 2.0913  
Epoch 3500, Loss 2.0907  
Epoch 4000, Loss 2.0907  
Epoch 4500, Loss 2.0907  
Final parameters: w2=0.8019, w1=9.1034, b=9.7710

Training with learning rate: 0.001  
Epoch 0, Loss 183.7949  
Epoch 500, Loss 157.5642  
Epoch 1000, Loss 135.0377  
Epoch 1500, Loss 115.7226  
Epoch 2000, Loss 99.2051  
Epoch 2500, Loss 85.1278  
Epoch 3000, Loss 73.1675  
Epoch 3500, Loss 63.0144  
Epoch 4000, Loss 54.3578  
Epoch 4500, Loss 46.8837  
Final parameters: w2=3.0189, w1=4.6637, b=4.2439

Training with learning rate: 0.0001  
Epoch 0, Loss 183.7949  
Epoch 500, Loss 180.9931  
Epoch 1000, Loss 178.2315  
Epoch 1500, Loss 175.5074  
Epoch 2000, Loss 172.8189  
Epoch 2500, Loss 170.1641  
Epoch 3000, Loss 167.5414  
Epoch 3500, Loss 164.9497  
Epoch 4000, Loss 162.3881  
Epoch 4500, Loss 159.8559  
Final parameters: w2=0.5891, w1=0.5965, b=0.4937



Best learning rate: 0.1

Best final loss: 2.0907

```
[10]: ## Problem 1
      ## 1c

      import torch
      import torch.optim as optim
      import matplotlib.pyplot as plt

      # Training data
      t_c = [0.5, 14.0, 15.0, 28.0, 11.0, 8.0, 3.0, -4.0, 6.0, 13.0, 21.0]
      t_u = [35.7, 55.9, 58.2, 81.9, 56.3, 48.9, 33.9, 21.8, 48.4, 60.4, 68.4]
      t_c = torch.tensor(t_c, dtype=torch.float)
      t_u = torch.tensor(t_u, dtype=torch.float)

      # Normalize input
      t_un = (t_u - torch.mean(t_u)) / torch.std(t_u)

      # Linear model
      def linear_model(t_u, w, b):
          return w * t_u + b
```

```

# Nonlinear model
def nonlinear_model(t_u, w2, w1, b):
    return w2 * t_u**2 + w1 * t_u + b

def loss_fn(t_p, t_c):
    squared_diffs = (t_p - t_c)**2
    return squared_diffs.mean()

# Train linear model
linear_params = torch.tensor([0.1, 0.0], requires_grad=True)
linear_optimizer = optim.Adam([linear_params], lr=0.01)

for epoch in range(5000):
    t_p = linear_model(t_un, *linear_params)
    loss = loss_fn(t_p, t_c)
    linear_optimizer.zero_grad()
    loss.backward()
    linear_optimizer.step()

linear_final_loss = loss.item()

# Train nonlinear model
nonlinear_params = torch.tensor([0.1, 0.1, 0.0], requires_grad=True)
nonlinear_optimizer = optim.Adam([nonlinear_params], lr=0.01)

for epoch in range(5000):
    t_p = nonlinear_model(t_un, *nonlinear_params)
    loss = loss_fn(t_p, t_c)
    nonlinear_optimizer.zero_grad()
    loss.backward()
    nonlinear_optimizer.step()

nonlinear_final_loss = loss.item()

# Plotting
plt.figure(figsize=(10, 6))
t_u_range = torch.linspace(min(t_un), max(t_un), 100)

# Plot data points
plt.scatter(t_un.numpy(), t_c.numpy(), label='Data', color='blue')

# Plot linear model predictions
linear_predictions = linear_model(t_u_range, *linear_params.detach())
plt.plot(t_u_range.numpy(), linear_predictions.numpy(), 'r-', label=f'Linear_
↳ (Loss: {linear_final_loss:.2f})')

# Plot nonlinear model predictions

```

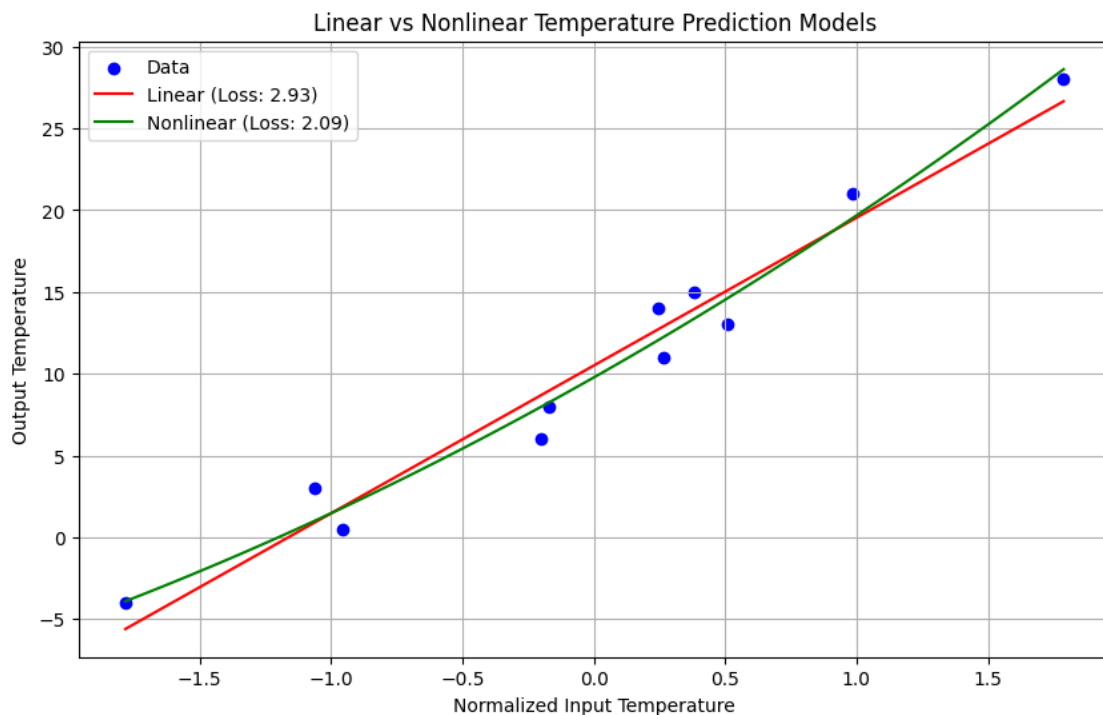
```

nonlinear_predictions = nonlinear_model(t_u_range, *nonlinear_params.detach())
plt.plot(t_u_range.numpy(), nonlinear_predictions.numpy(), 'g-', label=f'Nonlinear (Loss: {nonlinear_final_loss:.2f})')

plt.xlabel('Normalized Input Temperature')
plt.ylabel('Output Temperature')
plt.legend()
plt.title('Linear vs Nonlinear Temperature Prediction Models')
plt.grid(True)
plt.show()

print(f"\nLinear Model Final Loss: {linear_final_loss:.4f}")
print(f"Nonlinear Model Final Loss: {nonlinear_final_loss:.4f}")
print(f"Improvement: {((linear_final_loss - nonlinear_final_loss) / linear_final_loss * 100):.2f}%")

```



Linear Model Final Loss: 2.9276  
 Nonlinear Model Final Loss: 2.0907  
 Improvement: 28.59%

```

[17]: import torch
import torch.optim as optim
import pandas as pd

```

```

import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Load and preprocess data
data = pd.read_csv('assets/Housing.csv')
features = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking']
X = data[features]
y = data['price'] / 1e6 # Convert to millions for better scaling

# Split and scale
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Convert to tensors
X_train_tensor = torch.FloatTensor(X_train_scaled)
y_train_tensor = torch.FloatTensor(y_train.values)
X_test_tensor = torch.FloatTensor(X_test_scaled)
y_test_tensor = torch.FloatTensor(y_test.values)

class LinearRegression(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = torch.nn.Linear(5, 1)

    def forward(self, x):
        return self.linear(x)

# Training function with loss tracking
def train_and_plot(learning_rate):
    model = LinearRegression()
    criterion = torch.nn.MSELoss()
    optimizer = optim.SGD(model.parameters(), lr=learning_rate)

    train_losses = []
    epochs = []

    for epoch in range(5000):
        optimizer.zero_grad()
        outputs = model(X_train_tensor)
        loss = criterion(outputs, y_train_tensor.reshape(-1, 1))
        loss.backward()
        optimizer.step()

```



```

        if epoch % 500 == 0:
            epochs.append(epoch)
            train_losses.append(loss.item())
            print(f'Epoch {epoch}, Loss: {loss.item():.4f}')

    return epochs, train_losses, model

# Plot training curves for different learning rates
plt.figure(figsize=(10, 6))
learning_rates = [0.1, 0.01, 0.001, 0.0001]
best_loss = float('inf')
best_model = None

for lr in learning_rates:
    epochs, losses, model = train_and_plot(lr)
    plt.plot(epochs, losses, marker='o', label=f'Learning Rate = {lr}')
    if losses[-1] < best_loss:
        best_loss = losses[-1]
        best_model = model

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training Loss vs Epochs for Different Learning Rates')
plt.legend()
plt.grid(True)
plt.show()

# Plot actual vs predicted prices
with torch.no_grad():
    y_pred = best_model(X_test_tensor)

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred.numpy(), alpha=0.5)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel('Actual Price (millions)')
plt.ylabel('Predicted Price (millions)')
plt.title('Actual vs Predicted House Prices')
plt.grid(True)
plt.show()


```

```

Epoch 0, Loss: 23.6618
Epoch 500, Loss: 1.3500
Epoch 1000, Loss: 1.3500
Epoch 1500, Loss: 1.3500
Epoch 2000, Loss: 1.3500
Epoch 2500, Loss: 1.3500
Epoch 3000, Loss: 1.3500

```

Epoch 3500, Loss: 1.3500  
Epoch 4000, Loss: 1.3500  
Epoch 4500, Loss: 1.3500  
Epoch 0, Loss: 25.7232  
Epoch 500, Loss: 1.3500  
Epoch 1000, Loss: 1.3500  
Epoch 1500, Loss: 1.3500  
Epoch 2000, Loss: 1.3500  
Epoch 2500, Loss: 1.3500  
Epoch 3000, Loss: 1.3500  
Epoch 3500, Loss: 1.3500  
Epoch 4000, Loss: 1.3500  
Epoch 4500, Loss: 1.3500  
Epoch 0, Loss: 27.0583  
Epoch 500, Loss: 4.8386  
Epoch 1000, Loss: 1.8445  
Epoch 1500, Loss: 1.4244  
Epoch 2000, Loss: 1.3625  
Epoch 2500, Loss: 1.3524  
Epoch 3000, Loss: 1.3506  
Epoch 3500, Loss: 1.3502  
Epoch 4000, Loss: 1.3501  
Epoch 4500, Loss: 1.3500  
Epoch 0, Loss: 23.3274  
Epoch 500, Loss: 19.1591  
Epoch 1000, Loss: 15.8106  
Epoch 1500, Loss: 13.1122  
Epoch 2000, Loss: 10.9317  
Epoch 2500, Loss: 9.1656  
Epoch 3000, Loss: 7.7323  
Epoch 3500, Loss: 6.5671  
Epoch 4000, Loss: 5.6183  
Epoch 4500, Loss: 4.8448



```

[4]: ##Problem 2
    ## 2a 2b 2c

import torch
import torch.optim as optim
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Load and preprocess data
data = pd.read_csv('assets/Housing.csv')
features = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking']
X = data[features]
y = data['price']

# Split and scale
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
scaler_X = StandardScaler()
scaler_y = StandardScaler()

X_train_scaled = scaler_X.fit_transform(X_train)
X_test_scaled = scaler_X.transform(X_test)
y_train_scaled = scaler_y.fit_transform(y_train.values.reshape(-1, 1)).flatten()
y_test_scaled = scaler_y.transform(y_test.values.reshape(-1, 1)).flatten()

# Convert to tensors
X_train_tensor = torch.FloatTensor(X_train_scaled)
y_train_tensor = torch.FloatTensor(y_train_scaled)
X_test_tensor = torch.FloatTensor(X_test_scaled)
y_test_tensor = torch.FloatTensor(y_test_scaled)

class LinearRegression(torch.nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = torch.nn.Linear(5, 1) # 5 features -> 1 output

    def forward(self, x):
        return self.linear(x)

def train_model(lr):
    model = LinearRegression()
    criterion = torch.nn.MSELoss()
    optimizer = optim.SGD(model.parameters(), lr=lr)

```

```

train_losses = []
val_losses = []

print(f"\nTraining with learning rate: {lr}")
for epoch in range(5000):
    # Training
    model.train()
    outputs = model(X_train_tensor)
    loss = criterion(outputs, y_train_tensor.reshape(-1, 1))

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # Validation
    model.eval()
    with torch.no_grad():
        val_outputs = model(X_test_tensor)
        val_loss = criterion(val_outputs, y_test_tensor.reshape(-1, 1))

    if (epoch + 1) % 500 == 0:
        print(f'Epoch {epoch+1}:')
        print(f'Training Loss: {loss.item():.6f}')
        print(f'Validation Loss: {val_loss.item():.6f}')
        train_losses.append(loss.item())
        val_losses.append(val_loss.item())

    return model, train_losses, val_losses

# Train with different learning rates
learning_rates = [0.1, 0.01, 0.001, 0.0001]
all_models = []
all_train_losses = []
all_val_losses = []

for lr in learning_rates:
    model, train_losses, val_losses = train_model(lr)
    all_models.append(model)
    all_train_losses.append(train_losses)
    all_val_losses.append(val_losses)

# Find best model
best_model_idx = np.argmin([losses[-1] for losses in all_val_losses])
best_model = all_models[best_model_idx]
best_lr = learning_rates[best_model_idx]

print("\nBest Model Parameters:")

```

```

for name, param in best_model.named_parameters():
    if name == 'linear.weight':
        print("\nFeature weights:")
        for feature, weight in zip(features, param.data.numpy().flatten()):
            print(f"{feature}: {weight:.6f}")
    elif name == 'linear.bias':
        print(f"\nBias: {param.data.numpy()[0]:.6f}")

print(f"\nBest learning rate: {best_lr}")
print(f"Final validation loss: {all_val_losses[best_model_idx][-1]:.6f}")

# Plot training curves
plt.figure(figsize=(10, 6))
epochs = np.arange(500, 5001, 500)
for i, lr in enumerate(learning_rates):
    plt.plot(epochs, all_train_losses[i], label=f'Train (lr={lr})', marker='o')
    plt.plot(epochs, all_val_losses[i], label=f'Val (lr={lr})', marker='o',
             linestyle='--')

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss vs Epochs')
plt.legend()
plt.grid(True)
plt.show()

```

```

Training with learning rate: 0.1
Epoch 500:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 1000:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 1500:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 2000:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 2500:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 3000:
Training Loss: 0.437832
Validation Loss: 0.743571
Epoch 3500:
Training Loss: 0.437832

```

Validation Loss: 0.743571  
Epoch 4000:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 4500:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 5000:  
Training Loss: 0.437832  
Validation Loss: 0.743571

Training with learning rate: 0.01

Epoch 500:  
Training Loss: 0.437834  
Validation Loss: 0.743727  
Epoch 1000:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 1500:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 2000:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 2500:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 3000:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 3500:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 4000:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 4500:  
Training Loss: 0.437832  
Validation Loss: 0.743571  
Epoch 5000:  
Training Loss: 0.437832  
Validation Loss: 0.743571

Training with learning rate: 0.001

Epoch 500:  
Training Loss: 0.507405  
Validation Loss: 0.860031  
Epoch 1000:

Training Loss: 0.454805  
Validation Loss: 0.761232  
Epoch 1500:  
Training Loss: 0.442914  
Validation Loss: 0.746025  
Epoch 2000:  
Training Loss: 0.439387  
Validation Loss: 0.743386  
Epoch 2500:  
Training Loss: 0.438311  
Validation Loss: 0.743124  
Epoch 3000:  
Training Loss: 0.437980  
Validation Loss: 0.743258  
Epoch 3500:  
Training Loss: 0.437878  
Validation Loss: 0.743394  
Epoch 4000:  
Training Loss: 0.437846  
Validation Loss: 0.743481  
Epoch 4500:  
Training Loss: 0.437837  
Validation Loss: 0.743528  
Epoch 5000:  
Training Loss: 0.437834  
Validation Loss: 0.743551

Training with learning rate: 0.0001

Epoch 500:  
Training Loss: 1.267752  
Validation Loss: 1.776989  
Epoch 1000:  
Training Loss: 1.084056  
Validation Loss: 1.561130  
Epoch 1500:  
Training Loss: 0.944256  
Validation Loss: 1.394576  
Epoch 2000:  
Training Loss: 0.837147  
Validation Loss: 1.265131  
Epoch 2500:  
Training Loss: 0.754547  
Validation Loss: 1.163834  
Epoch 3000:  
Training Loss: 0.690446  
Validation Loss: 1.084043  
Epoch 3500:  
Training Loss: 0.640401




Validation Loss: 1.020803  
Epoch 4000:  
Training Loss: 0.601103  
Validation Loss: 0.970388  
Epoch 4500:  
Training Loss: 0.570075  
Validation Loss: 0.929978  
Epoch 5000:  
Training Loss: 0.545447  
Validation Loss: 0.897418

Best Model Parameters:

Feature weights:  
area: 0.387517  
bedrooms: 0.065551  
bathrooms: 0.321314  
stories: 0.241037  
parking: 0.163968

Bias: 0.000001

Best learning rate: 0.001  
Final validation loss: 0.743551



```

[7]: ##PROBLEM 3

import torch
import torch.optim as optim
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import matplotlib.pyplot as plt

# Load data
data = pd.read_csv('assets/Housing.csv')

# Convert categorical variables to numeric
categorical_columns = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
                       'airconditioning', 'prefarea', 'furnishingstatus']
label_encoders = {}

for column in categorical_columns:
    label_encoders[column] = LabelEncoder()
    data[column] = label_encoders[column].fit_transform(data[column])

# Prepare features and target
X = data.drop('price', axis=1)
y = data['price']

# Split and scale
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
scaler_X = StandardScaler()
scaler_y = StandardScaler()

X_train_scaled = scaler_X.fit_transform(X_train)
X_test_scaled = scaler_X.transform(X_test)
y_train_scaled = scaler_y.fit_transform(y_train.values.reshape(-1, 1)).flatten()
y_test_scaled = scaler_y.transform(y_test.values.reshape(-1, 1)).flatten()

# Convert to tensors
X_train_tensor = torch.FloatTensor(X_train_scaled)
y_train_tensor = torch.FloatTensor(y_train_scaled)
X_test_tensor = torch.FloatTensor(X_test_scaled)
y_test_tensor = torch.FloatTensor(y_test_scaled)

class LinearRegression(torch.nn.Module):
    def __init__(self, input_size):
        super().__init__()
        self.linear = torch.nn.Linear(input_size, 1)

```

```

def forward(self, x):
    return self.linear(x)

def train_model(lr):
    model = LinearRegression(X_train.shape[1])
    criterion = torch.nn.MSELoss()
    optimizer = optim.SGD(model.parameters(), lr=lr)

    train_losses = []
    val_losses = []

    print(f"\nTraining with learning rate: {lr}")
    for epoch in range(5000):
        # Training
        model.train()
        outputs = model(X_train_tensor)
        loss = criterion(outputs, y_train_tensor.reshape(-1, 1))

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Validation
        model.eval()
        with torch.no_grad():
            val_outputs = model(X_test_tensor)
            val_loss = criterion(val_outputs, y_test_tensor.reshape(-1, 1))

        if (epoch + 1) % 500 == 0:
            print(f'Epoch {epoch+1}:')
            print(f'Training Loss: {loss.item():.6f}')
            print(f'Validation Loss: {val_loss.item():.6f}')
            train_losses.append(loss.item())
            val_losses.append(val_loss.item())

    return model, train_losses, val_losses

# Train with different learning rates
learning_rates = [0.1, 0.01, 0.001, 0.0001]
all_models = []
all_train_losses = []
all_val_losses = []

for lr in learning_rates:
    model, train_losses, val_losses = train_model(lr)
    all_models.append(model)

```

```

all_train_losses.append(train_losses)
all_val_losses.append(val_losses)

# Find best model
best_model_idx = np.argmin([losses[-1] for losses in all_val_losses])
best_model = all_models[best_model_idx]
best_lr = learning_rates[best_model_idx]

print("\nBest Model Parameters:")
for name, param in best_model.named_parameters():
    if name == 'linear.weight':
        print("\nFeature weights:")
        for feature, weight in zip(X_train.columns, param.data.numpy().
            ↪flatten()):
            print(f"{feature}: {weight:.6f}")
    elif name == 'linear.bias':
        print(f"\nBias: {param.data.numpy()[0]:.6f}")

print(f"\nBest learning rate: {best_lr}")
print(f"Final validation loss: {all_val_losses[best_model_idx][-1]:.6f}")

# Plot training curves
plt.figure(figsize=(12, 6))
epochs = np.arange(500, 5001, 500)
for i, lr in enumerate(learning_rates):
    plt.plot(epochs, all_train_losses[i], label=f'Train (lr={lr})', marker='o')
    plt.plot(epochs, all_val_losses[i], label=f'Val (lr={lr})', marker='o',
        ↪linestyle='--')

plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss vs Epochs')
plt.legend()
plt.grid(True)
plt.show()

# Plot feature importance
plt.figure(figsize=(12, 6))
weights = best_model.linear.weight.data.numpy().flatten()
feature_importance = pd.DataFrame({
    'Feature': X_train.columns,
    'Weight': np.abs(weights)
}).sort_values('Weight', ascending=True)

plt.barh(feature_importance['Feature'], feature_importance['Weight'])
plt.xlabel('Absolute Weight Value')
plt.ylabel('Features')

```

```
plt.title('Feature Importance in Housing Price Prediction')
plt.tight_layout()
plt.show()
```

Training with learning rate: 0.1

Epoch 500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 1000:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 1500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 2000:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 2500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 3000:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 3500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 4000:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 4500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 5000:

Training Loss: 0.314557

Validation Loss: 0.574611

Training with learning rate: 0.01

Epoch 500:

Training Loss: 0.314575

Validation Loss: 0.574551

Epoch 1000:

Training Loss: 0.314557

Validation Loss: 0.574610

Epoch 1500:

Training Loss: 0.314557

Validation Loss: 0.574611

Epoch 2000:

Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 2500:  
Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 3000:  
Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 3500:  
Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 4000:  
Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 4500:  
Training Loss: 0.314557  
Validation Loss: 0.574611  
Epoch 5000:  
Training Loss: 0.314557  
Validation Loss: 0.574611

Training with learning rate: 0.001

Epoch 500:  
Training Loss: 0.399782  
Validation Loss: 0.658702  
Epoch 1000:  
Training Loss: 0.342636  
Validation Loss: 0.584102  
Epoch 1500:  
Training Loss: 0.325162  
Validation Loss: 0.573324  
Epoch 2000:  
Training Loss: 0.318813  
Validation Loss: 0.572084  
Epoch 2500:  
Training Loss: 0.316341  
Validation Loss: 0.572587  
Epoch 3000:  
Training Loss: 0.315328  
Validation Loss: 0.573241  
Epoch 3500:  
Training Loss: 0.314897  
Validation Loss: 0.573739  
Epoch 4000:  
Training Loss: 0.314709  
Validation Loss: 0.574071  
Epoch 4500:  
Training Loss: 0.314626

Validation Loss: 0.574280  
Epoch 5000:  
Training Loss: 0.314588  
Validation Loss: 0.574409

Training with learning rate: 0.0001

Epoch 500:  
Training Loss: 0.880955  
Validation Loss: 1.466555  
Epoch 1000:  
Training Loss: 0.710533  
Validation Loss: 1.219520  
Epoch 1500:  
Training Loss: 0.598023  
Validation Loss: 1.048471  
Epoch 2000:  
Training Loss: 0.522442  
Validation Loss: 0.928042  
Epoch 2500:  
Training Loss: 0.470674  
Validation Loss: 0.841833  
Epoch 3000:  
Training Loss: 0.434458  
Validation Loss: 0.779115  
Epoch 3500:  
Training Loss: 0.408556  
Validation Loss: 0.732782  
Epoch 4000:  
Training Loss: 0.389608  
Validation Loss: 0.698062  
Epoch 4500:  
Training Loss: 0.375439  
Validation Loss: 0.671705  
Epoch 5000:  
Training Loss: 0.364620  
Validation Loss: 0.651461

Best Model Parameters:

Feature weights:  
area: 0.293081  
bedrooms: 0.036938  
bathrooms: 0.298370  
stories: 0.192993  
mainroad: 0.074730  
guestroom: 0.053281  
basement: 0.103163  
hotwaterheating: 0.086414

airconditioning: 0.208585  
 parking: 0.109573  
 prefarea: 0.152070  
 furnishingstatus: -0.090537

Bias: 0.000001

Best learning rate: 0.001

Final validation loss: 0.574409

